

comp3009_assignment

September 24, 2019

```
[1]: import pandas as pd
import numpy as np
from sklearn import tree
from sklearn.utils import resample

import cufflinks as cf
import ipywidgets as widgets
from ipywidgets import interact, interact_manual
from ipywidgets.embed import embed_minimal_html

cf.go_offline(connected = False)
```

```
[2]: df = pd.read_csv('data2019.student.csv')
```

```
[3]: percent_missing = df.isnull().sum() * 100 / len(df)
missing_value_df = pd.DataFrame({'column_name': df.columns,
                                'percent_missing': percent_missing})
missing_value_df[missing_value_df['percent_missing'] > 0]
```

```
[3]:
```

	column_name	percent_missing
Class	Class	9.090909
att3	att3	0.363636
att9	att9	0.454545
att13	att13	93.454545
att19	att19	94.000000
att25	att25	0.272727
att28	att28	0.545455

```
[4]: to_drop = ['ID', 'att13', 'att19']
df.drop(labels = to_drop, axis = 1, inplace = True)

df['att3'].fillna(value = df['att3'].mode()[0], inplace = True)
df['att9'].fillna(value = df['att9'].mode()[0], inplace = True)
```

```
[5]: # percent_missing = df.isnull().sum() * 100 / len(df)
# missing_value_df = pd.DataFrame({'column_name': df.columns,
#                                   'percent_missing': percent_missing})
#
# missing_value_df
```

```

[6]: df['att25'].iplot(kind = 'box',
                        title = 'Boxplot of Attribute 25',
                        xTitle = '',
                        yTitle = 'value')

[7]: df['att28'].iplot(kind = 'box',
                        title = 'Boxplot of Attribute 28',
                        xTitle = '',
                        yTitle = 'value')

[8]: # median
df['att25'].fillna(value = int(df['att25'].median()), inplace = True)
# mean
df['att28'].fillna(value = int(np.floor(df['att28'].mean()))), inplace = True)

[9]: percent_missing = df.isnull().sum() * 100 / len(df)
missing_value_df = pd.DataFrame({'column_name': df.columns,
                                  'percent_missing': percent_missing})
missing_value_df[missing_value_df['percent_missing'] > 0]

[9]:      column_name  percent_missing
Class          Class          9.090909

[10]: # for i in df.columns:
#      if df[i].nunique() <= 1:
#          print(i)

[11]: to_drop = ['att14', 'att17']
df.drop(labels = to_drop, axis = 1, inplace = True)

[12]: # df.shape

[13]: df = df.T.drop_duplicates().T

[14]: # df.shape

[15]: df.drop_duplicates(inplace = True)

[16]: # df.shape

[17]: # @interact
# def box_plots(attribute=list(df.columns)):
#     df[attribute].iplot(kind = 'box')

[18]: df['att20'].iplot(kind = 'box', title = 'att20 Boxplot (numeric)')

[19]: df['att12'].iplot(kind = 'box', title = 'att12 Boxplot (categorical)')

[20]: df = df.apply(pd.to_numeric, errors='ignore')

[21]: #df.dtypes

[22]: to_numeric = ['att18', 'att20', 'att21', 'att22', 'att25', 'att28']

# for i in to_numeric:

```

```

#     display(df[i].describe())

[23]: # @interact_manual
# def hist_plots(attribute=list(to_numeric)):
#     df[attribute].iplot(kind = 'hist',
#                          title = str(attribute) + ' Before Scaling',
#                          xTitle = 'value')

[24]: df['att21'].iplot(kind = 'hist', title = 'att21 Before Scaling', xTitle = '
    ↳value')

[25]: to_gauss = ['att18', 'att20', 'att21', 'att22']
to_ln = ['att25', 'att28']

[26]: from sklearn import preprocessing

scaler = preprocessing.StandardScaler()
scaled = pd.DataFrame(scaler.fit_transform(df.loc[:,to_gauss].values), columns =
    ↳to_gauss)
df.loc[:,to_gauss] = scaled.values

[27]: scaler = preprocessing.RobustScaler()
scaled = pd.DataFrame(scaler.fit_transform(df.loc[:,to_ln].values), columns =
    ↳to_ln)
df.loc[:,to_ln] = scaled.values

[28]: scaler = preprocessing.MinMaxScaler()
scaled = pd.DataFrame(scaler.fit_transform(df.loc[:,to_numeric].values),
    ↳columns = to_numeric)
df.loc[:,to_numeric] = scaled.values

[29]: # @interact_manual
# def hist_plots(attribute=list(to_numeric)):
#     df[attribute].iplot(kind = 'hist',
#                          title = str(attribute) + ' After Scaling',
#                          xTitle = 'value')

[30]: df['att21'].iplot(kind = 'hist', title = 'att21 After Scaling', xTitle = '
    ↳value')

[31]: to_categorical = ['att30', 'att29', 'att27', 'att26', 'att23', 'att16',
    ↳'att15', 'att12', 'att11',
    ↳'att10', 'att9', 'att7', 'att6', 'att5', 'att4', 'att3',
    ↳'att2', 'att1']
df = pd.get_dummies(df, columns = to_categorical)

[32]: #df.shape

[33]: df['Class'].value_counts()

[33]: 1.0    650
0.0    250

```

Name: Class, dtype: int64

```
[34]: minority_class = df[df.Class == 0]
      minority_upsampled = resample(minority_class, replace = True, n_samples = (650_
      ↪ 250))

[35]: df_train = df[:-100]
      df_train = pd.concat([df_train, minority_upsampled])
      df_train = df_train.reset_index(drop = True)
      y_train = df_train.loc[:, 'Class'].values
      X_train = df_train.iloc[:, 1:].values

      df_test = df[-100:]
      X_test_final = df_test.iloc[:, 1:].values

[36]: from sklearn.feature_selection import SelectKBest
      from sklearn.feature_selection import chi2

      Xx = df_train.iloc[:, 1:]
      yy = df_train.loc[:, 'Class']

      bestfeatures = SelectKBest(score_func = chi2, k = 10)
      fit = bestfeatures.fit(Xx, yy)
      dfscores = pd.DataFrame(fit.scores_)
      dfcolumns = pd.DataFrame(Xx.columns)

      featureScores = pd.concat([dfcolumns, dfscores], axis = 1)
      featureScores.columns = ['Specs', 'Score']
      featureScores = featureScores.set_index('Specs').sort_values(by = ['Score'],
      ↪ axis = 0, ascending = True)
      #print(featureScores.nlargest(35, 'Score'))%xdel

      featureScores.iplot(kind='barh', title = 'Univariate Feature Importance (K-Best_
      ↪from Chi^2)')

[37]: feat_list_1 = featureScores[-15:].index

[38]: from sklearn.ensemble import ExtraTreesClassifier

      model = ExtraTreesClassifier(n_estimators = 100)
      model.fit(Xx, yy)
      #print(model.feature_importances_)
      feat_importances = pd.Series(model.feature_importances_, index=Xx.columns)
      feat_importances = feat_importances.sort_values()

      feat_importances.iplot(kind='barh', title = 'Tree-Based Feature Importance')

[39]: feat_list_2 = feat_importances[-11:].index
```

```
[40]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.feature_selection import SelectFromModel

# lsvc = LinearSVC(C=0.01, penalty="l1", dual=False).fit(Xx, yy)
# model = SelectFromModel(lsvc, prefit=True)
# feat_list_3 = Xx.iloc[:,list(model.get_support(indices=True))].columns

lr = LogisticRegression(C = 0.000000001, penalty = 'l2', dual = False, solver = 'lbfgs').fit(Xx, yy)
model = SelectFromModel(lr, prefit=True)
feat_list_3 = Xx.iloc[:,list(model.get_support(indices=True))].columns
#feat_list_3

[41]: # common within all three methods - probably can be considered rather important
      # as a result
      # list(set(feat_list_1) & set(feat_list_2) & set(feat_list_3))

[42]: # every single attribute seen across all three methods
important_cols = list(set(list(feat_list_1) + list(feat_list_2) +
      list(feat_list_3)))

X_train = Xx.loc[:,important_cols].values
X_test = df_test.loc[:,important_cols].values

[43]: from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, f1_score, accuracy_score

X_test_final = df_test.loc[:,important_cols].values
df_test.loc[:,important_cols].to_csv('final_100_set.csv', index = False)
X = df_train.loc[:,important_cols]
y = df_train.loc[:, 'Class']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
      random_state=6346)

[44]: # X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
      # random_state=6346)
      # X_train, X_val, y_train, y_val = train_test_split(X_train, y_train,
      # test_size=0.2, random_state=4653)

[45]: # pd.concat([y_val, X_val], axis=1).to_csv('val.csv', index = True)
      # pd.concat([y_test, X_test], axis=1).to_csv('test.csv', index = True)
      # pd.concat([y_train, X_train], axis=1).to_csv('train.csv', index = True)

[46]: # import numpy as np
```

```

# def pandas2arff(df,filename,wekaname =
    →"pandasdata",cleanstringdata=True,cleannan=True):
#     """
#     converts the pandas dataframe to a weka compatible file
#     df: dataframe in pandas format
#     filename: the filename you want the weka compatible file to be in
#     wekaname: the name you want to give to the weka dataset (this will be
    →visible to you when you open it in Weka)
#     cleanstringdata: clean up data which may have spaces and replace with
    →"_", special characters etc which seem to annoy Weka.
#         To suppress this, set this to False
#     cleannan: replaces all nan values with "?" which is Weka's standard for
    →missing values.
#         To suppress this, set this to False
#     """
#     import re

#     def cleanstring(s):
#         if s!="?":
#             return re.sub('[^A-Za-z0-9]+', "_", str(s))
#         else:
#             return "?"

#     dfcopy = df #all cleaning operations get done on this copy

#     if cleannan!=False:
#         dfcopy = dfcopy.fillna(-999999999) #this is so that we can swap this
    →out for "?"
#         #this makes sure that certain numerical columns with missing values
    →don't get stuck with "object" type

#     f = open(filename,"w")
#     arffList = []
#     arffList.append("@relation " + wekaname + "\n")
#     #look at each column's dtype. If it's an "object", make it "nominal"
    →under Weka for now (can be changed in source for dates.. etc)
#     for i in range(df.shape[1]):
#         if dfcopy.dtypes[i]=='O' or (df.columns[i] in
    →["Class","CLASS","class"]):
#             if cleannan!=False:
#                 dfcopy.iloc[:,i] = dfcopy.iloc[:,i].
    →replace(to_replace=-999999999, value="?")
#             if cleanstringdata!=False:
#                 dfcopy.iloc[:,i] = dfcopy.iloc[:,i].apply(cleanstring)

```

```

#         _uniqueNominalVals = [str(_i) for _i in np.unique(dfcopy.iloc[:
    →, i])]
#         _uniqueNominalVals = ",".join(_uniqueNominalVals)
#         _uniqueNominalVals = _uniqueNominalVals.replace("[", "")
#         _uniqueNominalVals = _uniqueNominalVals.replace("]", "")
#         _uniqueValuesString = "{" + _uniqueNominalVals + "}"
#         arffList.append("@attribute " + df.columns[i] + "
    → _uniqueValuesString + "\n")
#     else:
#         arffList.append("@attribute " + df.columns[i] + " real\n")
#         #even if it is an integer, let's just deal with it as a real
    → number for now
#         arffList.append("@data\n")
#         for i in range(dfcopy.shape[0]):#instances
#             _instanceString = ""
#             for j in range(df.shape[1]):#features
#                 if dfcopy.dtypes[j]=='O':
#                     _instanceString+="\""+ str(dfcopy.iloc[i,j]) + "\""
#                 else:
#                     _instanceString+=str(dfcopy.iloc[i,j])
#                 if j!=dfcopy.shape[1]-1:#if it's not the last feature, add a
    → comma
#                     _instanceString+=", "
#             _instanceString+="\n"
#             if cleannan!=False:
#                 _instanceString = _instanceString.replace("-999999999.0", "?")
    → #for numeric missing values
#             _instanceString = _instanceString.replace("\"?\"", "?") #for
    → categorical missing values
#         arffList.append(_instanceString)
#         f.writelines(arffList)
#         f.close()
#         del dfcopy
#         return True

```

```

[47]: # pandas2arff(pd.concat([y_val, X_val], axis=1), 'val.arff', wekaname="val",
    → cleanstringdata=False, cleannan=False)
# pandas2arff(pd.concat([y_test, X_test], axis=1), 'test.arff',
    → wekaname="test", cleanstringdata=False, cleannan=False)
# pandas2arff(pd.concat([y_train, X_train], axis=1), 'train.arff',
    → wekaname="train", cleanstringdata=False, cleannan=False)

```

```

[48]: from sklearn.model_selection import cross_val_score, cross_val_predict

from sklearn.cluster import KMeans
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression

```

```

from sklearn import svm
from sklearn.naive_bayes import GaussianNB, MultinomialNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier

results = pd.DataFrame()
final_proportions = pd.DataFrame()

all_models = dict([
    ('kmeans', KMeans(n_clusters = 2)),
    ('rForest', RandomForestClassifier(n_estimators=100, max_depth=2)),
    ('logReg', LogisticRegression(solver = 'lbfgs', max_iter = 2500)),
    ('svm', svm.SVC(gamma='scale', decision_function_shape='ovo')),
    ('nBayes', MultinomialNB(alpha=2.5)),
    ('adaBoost', AdaBoostClassifier(n_estimators=100)),
    ('knearest', KNeighborsClassifier(n_neighbors=73, p=1, weights="distance"))
])

for i in all_models.keys():
    results[i + '_acc'] = cross_val_score(all_models[i], X_train, y_train,
    →cv=10, scoring = 'accuracy')
    results[i + '_f1'] = cross_val_score(all_models[i], X_train, y_train,
    →cv=10, scoring = 'f1')

    all_models[i].fit(X_train, y_train)
    pred_final = all_models[i].predict(X_test_final)
    zeros = np.bincount(pred_final.astype(int))[0]
    ones = np.bincount(pred_final.astype(int))[1]

    final_proportions[i] = [zeros, ones, (abs(ones - zeros) / 100), (1 -
    →abs(ones - zeros) / 100)]

```

[49]:

```

# from sklearn.pipeline import Pipeline
# from sklearn.model_selection import GridSearchCV

# pipe = Pipeline([('classifier' , svm.SVC())])

# param_grid = [
# #     {'classifier' : [LogisticRegression()],
# #     'classifier__penalty' : ['l1', 'l2'],
# #     'classifier__C' : np.logspace(-4, 4, 20),
# #     'classifier__solver' : ['liblinear']},
# #     {'classifier' : [RandomForestClassifier()],
# #     'classifier__n_estimators' : list(range(10,101,10)),
# #     'classifier__max_features' : list(range(6,32,5))},

```



```
# #      {'classifier' : [MultinomialNB()],
# #      'classifier__alpha' : [0.2,0.5,0.6,1,1.5,2.
→5,5,10,20,50,75,90,130,250,500]},
#      {'classifier' : [svm.SVC()],
#      'classifier__kernel' : ['linear', 'poly', 'rbf', 'sigmoid'],
#      'classifier__C' : [1,2,5,10,100,50,90],
#      'classifier__gamma' : ['scale', 1, 2, 3, 5, 0.5, 100, 25000],
#      'classifier__decision_function_shape' : ['ovr', 'ovo']}

# ]

# clf = GridSearchCV(pipe, param_grid = param_grid, cv = 5, verbose=True,
→n_jobs=-1)

# best_clf = clf.fit(X_train, y_train)

# best_clf.best_estimator_
# best_clf.best_params_
```

```
[50]: results.iplot(kind = 'box',
                    colors = ['#172144', '#172144', '#306E46', '#306E46', '#5D9E39',
→'#5D9E39',
                                '#6E0D09', '#6E0D09', '#1D3168', '#1D3168', '#286263',
→'#286263',
                                '#DB7C26', '#DB7C26'],
                    yrange=[0,1],
                    title = 'Classification Comparison',
                    layout_update=dict([('yaxis', {'nticks':11}),
                                        ('margin', {'b':100})]))
#                                , asPlot = True) ## uncomment (and remove bracket) to display
→in browser
```

```
[51]: final_proportions.index = ['n_zeros', 'n_ones', 'difference', 'correct']
final_proportions
```

```
[51]:
```

	kmeans	rForest	logReg	svm	nBayes	adaBoost	knearest
n_zeros	54.00	52.00	49.00	45.0	47.00	42.00	51.00
n_ones	46.00	48.00	51.00	55.0	53.00	58.00	49.00
difference	0.08	0.04	0.02	0.1	0.06	0.16	0.02
correct	0.92	0.96	0.98	0.9	0.94	0.84	0.98

```
[52]: df_predictions = pd.DataFrame()
df_predictions['ID'] = [i for i in range(1001, 1101)]

all_models['knearest'].fit(X_train, y_train)
df_predictions['Predict 1'] = all_models['knearest'].predict(X_test_final)
```

```
all_models['logReg'].fit(X_train, y_train)
df_predictions['Predict 2'] = all_models['logReg'].predict(X_test_final)

df_predictions = df_predictions.astype(int)
```

```
[53]: #df_predictions.to_csv('predict_actual.csv', index = False)
```

```
[54]: df_predictions[0:10]
```

```
[54]:
```

	ID	Predict 1	Predict 2
0	1001	0	0
1	1002	0	0
2	1003	0	0
3	1004	0	0
4	1005	0	0
5	1006	1	1
6	1007	1	1
7	1008	0	0
8	1009	1	1
9	1010	0	0